

# financial\_risk\_eda

July 29, 2025

## 1 CS 6220 Final Project: Predicting Firms' Financial Risk

Group Members: Jiajun Fang, Yini Li

```
[1]: # Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
import warnings
warnings.filterwarnings('ignore')
```

## 2 Load Parquet Files

```
[2]: # Step 1: Import required libraries
import pandas as pd
import os
```

```
[3]: # Step 2: Check if the Parquet files exist in the current directory
train_file = "20231124_Financial_Risk_Project_train (1).parquet" # We will only
    ↪ use this file for final project because test_file has no FinancialRisk
    ↪ measurement
# test_file = "20231124_Financial_Risk_Project_test_public (1).parquet"

print("Train file exists:", os.path.exists(train_file))
# print("Test file exists:", os.path.exists(test_file))
```

Train file exists: True

```
[4]: # Step 3: Load the parquet files
train_df = pd.read_parquet(train_file)
# test_df = pd.read_parquet(test_file)

print("Train shape:", train_df.shape)
# print("Test shape:", test_df.shape)

train_df.head()
```

Train shape: (415, 73)

```
[4]:                                     url \
ticker
FICO    https://seekingalpha.com/article/4649507-fair-...
RXRX    https://seekingalpha.com/earnings/earnings-cal...
BLDR    https://seekingalpha.com/article/4645938-build...
JKHY    https://seekingalpha.com/article/4649242-jack-...
OC      https://seekingalpha.com/article/4643409-owens...
```

```
                                     call_transcript    VWAP \
ticker
FICO    Fair Isaac Corporation (FICO) Q4 2023 Earnings...  700.83
RXRX    Earnings Call Transcripts | Seeking Alpha\n\n...    6.67
BLDR    Builders FirstSource, Inc. (BLDR) Q3 2023 Earn...  88.34
JKHY    Jack Henry & Associates, Inc. (JKHY) Q1 2024 E... 150.47
OC      Owens Corning (OC) Q3 2023 Earnings Call Trans...  95.47
```

```
exchangeCountry    securityType    CIK \
ticker
FICO              USA    Common or ordinary  0000814547
RXRX              USA    Common or ordinary  0001601830
BLDR              USA    Common or ordinary  0001316835
JKHY              USA    Common or ordinary  0000779152
OC                USA    Common or ordinary  0001370946
```

```
name    securityID    incorporationCountry \
ticker
FICO    Fair Isaac Corporation    138240101    USA
RXRX    Recursion Pharmaceuticals Inc    384740101    USA
BLDR    Builders FirstSource Inc    1630360101    USA
JKHY    Henry (Jack) & Associates Inc    118110101    USA
OC      Owens Corning    82140601    USA
```

```
exchangeName    ... Working Capital \
ticker    ...
FICO    New York Stock Exchange    ...    5.063632
RXRX    Nasdaq Stock Market    ...    2.528944
BLDR    New York Stock Exchange    ...    11.272943
```

JKHY	Nasdaq Stock Market	...	0.843418
OC	New York Stock Exchange	...	15.414013

	businessDescription	close	\
ticker			
FICO	Fair Isaac Corporation develops analytic, soft...	598.58	
RXXR	Recursion Pharmaceuticals, Inc. operates as a ...	7.71	
BLDR	Builders FirstSource, Inc., together with its ...	64.88	
JKHY	Jack Henry & Associates, Inc., a financial tec...	175.56	
OC	Owens Corning engages in manufacture and sale ...	85.30	

	dividendFactor	fiscalDint	floatShares	outstandingShares	\
ticker					
FICO	1.000000	20221231	24272216.0	25154000.0	
RXXR	1.000000	20221231	117933131.0	174072906.0	
BLDR	1.000000	20221231	146461116.0	148994000.0	
JKHY	0.989936	20221231	72430647.0	72910225.0	
OC	0.980838	20221231	94098317.0	94700000.0	

	shortInterestFloat	FinancialSector	FinancialRisk
ticker			
FICO	0.033322	0	0.0
RXXR	0.130693	0	0.0
BLDR	0.040739	0	0.0
JKHY	0.030052	0	0.0
OC	0.038574	0	0.0

[5 rows x 73 columns]

```
[5]: # Use the already loaded parquet data
df = train_df.copy()
print(f"Dataset shape: {df.shape}")
df.head()
```

Dataset shape: (415, 73)

```
[5]:                                     url \
ticker
FICO  https://seekingalpha.com/article/4649507-fair-...
RXXR  https://seekingalpha.com/earnings/earnings-cal...
BLDR  https://seekingalpha.com/article/4645938-build...
JKHY  https://seekingalpha.com/article/4649242-jack-...
OC    https://seekingalpha.com/article/4643409-owens...
```

	call_transcript	VWAP	\
ticker			
FICO	Fair Isaac Corporation (FICO) Q4 2023 Earnings...	700.83	
RXXR	Earnings Call Transcripts   Seeking Alpha\n\n\n...	6.67	

BLDR	Builders FirstSource, Inc. (BLDR) Q3 2023 Earn...	88.34
JKHY	Jack Henry & Associates, Inc. (JKHY) Q1 2024 E...	150.47
OC	Owens Corning (OC) Q3 2023 Earnings Call Trans...	95.47

	exchangeCountry	securityType	CIK \
ticker			
FICO	USA	Common or ordinary	0000814547
RXXRX	USA	Common or ordinary	0001601830
BLDR	USA	Common or ordinary	0001316835
JKHY	USA	Common or ordinary	0000779152
OC	USA	Common or ordinary	0001370946

	name	securityID	incorporationCountry \
ticker			
FICO	Fair Isaac Corporation	138240101	USA
RXXRX	Recursion Pharmaceuticals Inc	384740101	USA
BLDR	Builders FirstSource Inc	1630360101	USA
JKHY	Henry (Jack) & Associates Inc	118110101	USA
OC	Owens Corning	82140601	USA

	exchangeName	... Working Capital \
ticker		
FICO	New York Stock Exchange	5.063632
RXXRX	Nasdaq Stock Market	2.528944
BLDR	New York Stock Exchange	11.272943
JKHY	Nasdaq Stock Market	0.843418
OC	New York Stock Exchange	15.414013

	businessDescription	close \
ticker		
FICO	Fair Isaac Corporation develops analytic, soft...	598.58
RXXRX	Recursion Pharmaceuticals, Inc. operates as a ...	7.71
BLDR	Builders FirstSource, Inc., together with its ...	64.88
JKHY	Jack Henry & Associates, Inc., a financial tec...	175.56
OC	Owens Corning engages in manufacture and sale ...	85.30

	dividendFactor	fiscalDint	floatShares	outstandingShares \
ticker				
FICO	1.000000	20221231	24272216.0	25154000.0
RXXRX	1.000000	20221231	117933131.0	174072906.0
BLDR	1.000000	20221231	146461116.0	148994000.0
JKHY	0.989936	20221231	72430647.0	72910225.0
OC	0.980838	20221231	94098317.0	94700000.0

	shortInterestFloat	FinancialSector	FinancialRisk
ticker			
FICO	0.033322	0	0.0

RXRX	0.130693	0	0.0
BLDR	0.040739	0	0.0
JKHY	0.030052	0	0.0
OC	0.038574	0	0.0

[5 rows x 73 columns]

## 2.1 Data Cleaning

```
[6]: # Drop irrelevant features
irrelevant_cols = ['url', 'exchangeCountry', 'CIK', 'securityID',
                  'incorporationCountry', 'exchangeName', 'exchangeID',
                  'businessDescription', 'FinancialSector']

# Drop FinancialSector because the variable describes which company is a
# finance company and we should not know this for the purpose of this project
df.drop(columns=irrelevant_cols, inplace=True, errors='ignore')

# Drop FinancialSector variable?

# Drop columns with more than 20% missing values
missing_threshold = 0.2
df.dropna(axis=1, thresh=(1 - missing_threshold) * len(df), inplace=True)

# Separate numerical and categorical columns
num_cols = df.select_dtypes(include=['float64', 'int64']).columns.tolist()
cat_cols = df.select_dtypes(include=['object']).columns.tolist()
```

## 2.2 Exploratory Data Analysis (EDA)

```
[7]: df.describe()
```

```
[7]:
```

	VWAP	Accrual Ratio	Assets	B/P	CF/P \
count	415.000000	415.000000	415.000000	415.000000	415.000000
mean	116.474308	-0.018919	200.850240	0.527509	0.112481
std	216.885791	0.176063	408.857675	0.467866	0.195024
min	0.199900	-3.202525	1.228840	-0.636664	-0.775506
25%	24.995000	-0.030362	37.962225	0.170290	0.035302
50%	57.950000	-0.005417	78.088405	0.454349	0.079560
75%	130.317500	0.018226	183.633250	0.763030	0.166526
max	2648.960000	0.261432	4042.049341	3.750665	1.727142

	Capital Expenditure	Cash	Debt/Equity	Depreciation \
count	415.000000	415.000000	397.000000	415.000000
mean	2.303081	18.136927	1.914057	1.930943
std	4.525536	70.680449	7.141619	3.588728
min	0.000000	0.000000	0.000000	0.000000

25%	0.123280	2.487505	0.337896	0.123215
50%	0.701300	5.497844	0.738463	0.810170
75%	2.418535	13.397640	1.463662	2.658760
max	57.515690	1309.060835	116.069725	54.195690

	Dividend	...	Short Term Debt	TEV	Working Capital	\
count	415.000000	...	415.000000	399.000000	415.000000	
mean	1.604028	...	9.638088	132.510298	5.553125	
std	2.162034	...	48.195878	215.511729	17.098716	
min	0.000000	...	0.000000	-63.167179	-105.669178	
25%	0.000000	...	0.180620	37.830784	0.000000	
50%	1.036192	...	1.153693	75.609645	0.000000	
75%	2.111366	...	4.958480	153.935241	6.529998	
max	19.611487	...	882.360527	2924.920173	189.206644	

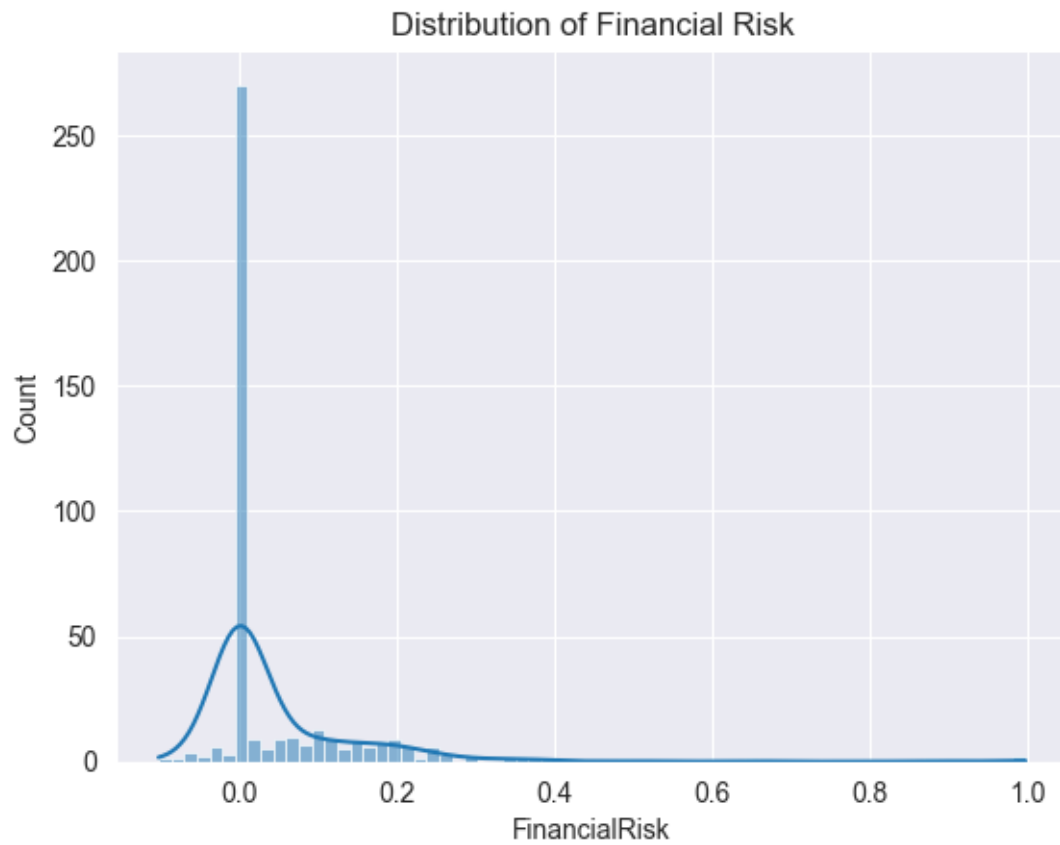
	close	dividendFactor	fiscalDint	floatShares	\
count	415.000000	415.000000	4.150000e+02	4.150000e+02	
mean	110.817879	0.978670	2.022126e+07	5.162328e+08	
std	190.034596	0.025503	6.180495e+02	1.397186e+09	
min	1.830000	0.863362	2.022103e+07	1.305533e+07	
25%	27.245000	0.968785	2.022123e+07	7.171691e+07	
50%	61.350000	0.985036	2.022123e+07	1.518589e+08	
75%	130.060000	1.000000	2.022123e+07	3.686314e+08	
max	2460.840000	1.000000	2.023010e+07	1.592372e+10	

	outstandingShares	shortInterestFloat	FinancialRisk
count	4.150000e+02	414.000000	415.000000
mean	5.584309e+08	0.039076	0.049077
std	1.497614e+09	0.039058	0.117747
min	1.365940e+07	0.002825	-0.102909
25%	7.712301e+07	0.015174	0.000000
50%	1.665872e+08	0.026484	0.000000
75%	3.947026e+08	0.049398	0.060544
max	1.594342e+10	0.287963	0.998330


[8 rows x 56 columns]

```
[8]: # Check distribution of target
sns.histplot(df['FinancialRisk'], kde=True)
plt.title('Distribution of Financial Risk')
plt.show()

# Log-transform target
df['FinancialRisk_log'] = np.log1p(df['FinancialRisk'])
```



```
[9]: # Correlation matrix
plt.figure(figsize=(12, 10))
corr = df[num_cols].corr()
sns.heatmap(corr, cmap='coolwarm', center=0)
plt.title('Correlation Matrix')
plt.show()
```



```
[10]: # Distribution of some variables with the highest correlation score
```

```
# Get correlation with the outcome variable
target_corr = corr['FinancialRisk'].drop('FinancialRisk')

# Select features with |correlation| > 0.2
selected = target_corr[abs(target_corr) > 0.2]

# Print feature names and their correlation scores
print("Selected Features and Correlation Scores:\n")
for feature, score in selected.items():
    print(f"{feature}: {score:.3f}")
```

Selected Features and Correlation Scores:

Assets: 0.509



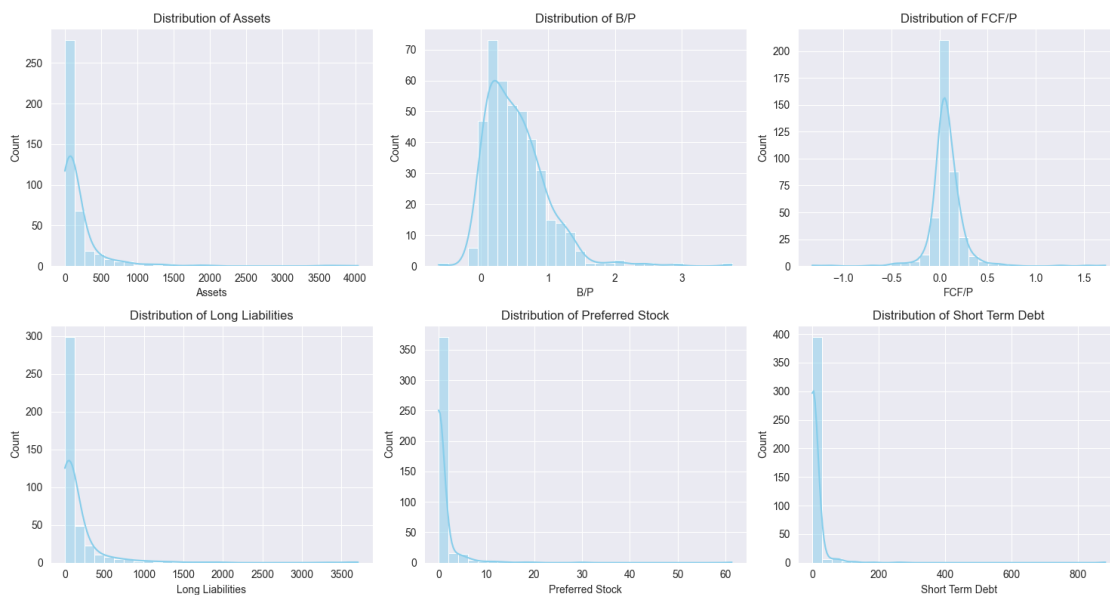
B/P: 0.289  
 FCF/P: 0.214  
 Long Liabilities: 0.514  
 Preferred Stock: 0.495  
 Short Term Debt: 0.264

```
[11]: # Set up the plot grid
# Plot the variables distribution (histogram) using for loops
num_feats = len(selected)
cols = 3
rows = (num_feats + cols - 1) // cols

plt.figure(figsize=(5 * cols, 4 * rows))

for i, col in enumerate(selected.index, 1): # use selected.index instead of selected
    plt.subplot(rows, cols, i)
    sns.histplot(df[col], kde=True, bins=30, color='skyblue')
    plt.title(f'Distribution of {col}')

plt.tight_layout()
plt.show()
```



```
[12]: # Scatter plots between the highest correlated variable with outcome variable
      ↪ (FinancialRisk

# Scatter plots between selected features and FinancialRisk
```

```

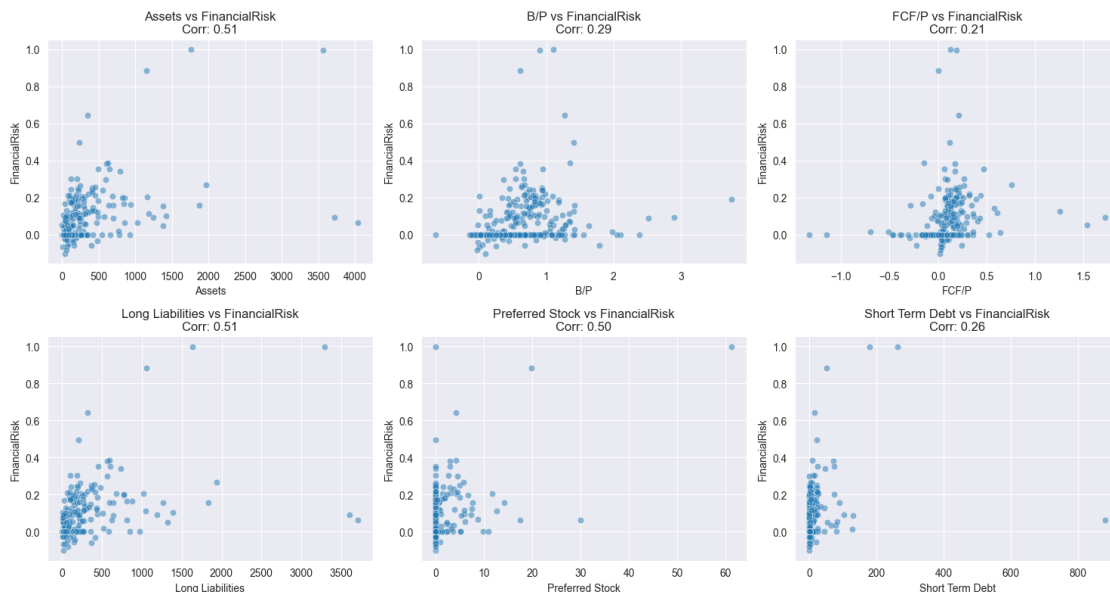
num_feats = len(selected)
cols = 3
rows = (num_feats + cols - 1) // cols

plt.figure(figsize=(5 * cols, 4 * rows))

for i, col in enumerate(selected.index, 1):
    plt.subplot(rows, cols, i)
    sns.scatterplot(data=df, x=col, y='FinancialRisk', alpha=0.5)
    plt.title(f'{col} vs FinancialRisk\nCorr: {selected[col]:.2f}')
    plt.xlabel(col)
    plt.ylabel('FinancialRisk')

plt.tight_layout()
plt.show()

```



## 2.3 Baseline Model: Linear Regression

```

[13]: # Baseline Model: Linear Regression
selected_features = ['Assets', 'Cash', 'Debt/Equity', 'B/P', 'SG&A/Sales',
                    ↪ 'Long Liabilities']
print('Columns in DataFrame:', df.columns.tolist())
missing_features = [col for col in selected_features if col not in df.columns]
if missing_features:
    print('The following features are missing from the data:', missing_features)
else:
    print('All selected features are present in the data.')

```

```

if 'FinancialRisk_log' not in df.columns:
    print('FinancialRisk_log column is missing. Please make sure the previous_
    ↪cell has been executed.')
else:
    print('FinancialRisk_log column exists.')
print('Missing value count for selected features:')
print(df[selected_features].isnull().sum())
# Convert selected features to numeric type
for col in selected_features:
    df[col] = pd.to_numeric(df[col], errors='coerce')
# Remove all rows with missing values in selected features or target
df_clean = df.dropna(subset=selected_features + ['FinancialRisk_log']).copy()
print('Shape after dropping missing:', df_clean.shape)
# Continue with modeling
X = df_clean[selected_features]
y = df_clean['FinancialRisk_log']
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, X.columns)
    ])
from sklearn.linear_model import LinearRegression
model = Pipeline(steps=[('preprocessor', preprocessor),
    ('regressor', LinearRegression())])
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
from sklearn.metrics import mean_squared_error, r2_score
r2 = r2_score(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred) ** 0.5
print(f'R2 Score: {r2:.2f}')
print(f'RMSE: {rmse:.2f}')

```

Columns in DataFrame: ['call\_transcript', 'VWAP', 'securityType', 'name', 'Accrual Ratio', 'Assets', 'B/P', 'CF/P', 'Capital Expenditure', 'Cash', 'Debt/Equity', 'Depreciation', 'Dividend', 'E/P', 'EBIT', 'EBIT/P', 'EBIT/TEV', 'Earnings', 'Earnings Growth (1Y)', 'Equity', 'FCF', 'FCF/P', 'Income Tax', 'Interest Expense', 'Long Liabilities', 'Long Term Debt', 'Market Cap',

```
'Minority Interest', 'Operating Cash Flow', 'Operating Expense', 'Operating  
Income', 'Operating Income Before Depreciation', 'Operating Margin', 'Preferred  
Stock', 'Profit Margin', 'R&D', 'R&D/Sales', 'ROA', 'ROE', 'S/P', 'SG&A',  
'SG&A/Sales', 'Sales', 'Sales Growth (1Y)', 'Sales Growth (2Y)', 'Sales Growth  
(3Y)', 'Sales Growth (4Y)', 'Sales Growth (5Y)', 'Sales Variability', 'Short  
Term Debt', 'TEV', 'Working Capital', 'close', 'dividendFactor', 'fiscalDint',  
'floatShares', 'outstandingShares', 'shortInterestFloat', 'FinancialRisk',  
'FinancialRisk_log']
```

All selected features are present in the data.

FinancialRisk\_log column exists.

Missing value count for selected features:

Assets	0
Cash	0
Debt/Equity	18
B/P	0
SG&A/Sales	7
Long Liabilities	0


dtype: int64

Shape after dropping missing: (390, 60)

R<sup>2</sup> Score: -0.08

RMSE: 0.08

```
[14]: # Test Co-linearity between the highly correlated variables  
  
# Correlation heatmap of selected features  
plt.figure(figsize=(10, 8))  
sns.heatmap(df[selected.index].corr(), annot=True, cmap='coolwarm', center=0)  
plt.title("Collinearity Heatmap Among Selected Features")  
plt.show()
```



```
[15]: # Further examines multicollinearity
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.tools.tools import add_constant
import pandas as pd

X = df[selected.index].dropna()
X = add_constant(X)

vif_data = pd.DataFrame()
vif_data['Feature'] = X.columns
vif_data['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.
    ↪shape[1])]

print(vif_data)
# Baseline model might have low performance because Asset and Long Liabilities_
    ↪are strongly correlated with each other
```

	Feature	VIF
0	const	2.703747
1	Assets	143.285350
2	B/P	1.101047
3	FCF/P	1.182681
4	Long Liabilities	137.109055
5	Preferred Stock	1.878595
6	Short Term Debt	1.873132

```
[16]: # Drop Assets or Long Liabilities (drop asset) because it is highly correlated
      ↪with each other
      reduced_features = selected.index.drop('Assets')
```

## 2.4 Modeling 2: Regularization Techniques

Using Lasso, Ridge, and elastic net + stepwise selection

Try the regularization models with all the features

```
[17]: from sklearn.preprocessing import StandardScaler
      from sklearn.linear_model import LassoCV, RidgeCV, ElasticNetCV

      full_features = df.drop(columns=['FinancialRisk',
      ↪'Assets', 'FinancialRisk_log']) # Drop target
      X = full_features.select_dtypes(include=[np.number]).dropna() # Keep only
      ↪numerical features
      y = df.loc[X.index, 'FinancialRisk']

      scaler = StandardScaler()
      X_scaled = scaler.fit_transform(X)
      X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
      ↪random_state=42)
```

```
[18]: # Train regularized models
      # Lasso
      lasso = LassoCV(cv=5, random_state=42)
      lasso.fit(X_train, y_train)

      # Ridge
      ridge = RidgeCV(alphas=[0.1, 1.0, 10.0], cv=5)
      ridge.fit(X_train, y_train)

      # ElasticNet
      elastic = ElasticNetCV(cv=5, random_state=42)
      elastic.fit(X_train, y_train)

      # Evaluation function
      def evaluate(model, name):
```

```

preds = model.predict(X_test)
r2 = r2_score(y_test, preds)
rmse = np.sqrt(mean_squared_error(y_test, preds))
print(f"{name} Performance:")
print(f"  R2: {r2:.3f}")
print(f"  RMSE: {rmse:.3f}")
print("-" * 30)
return r2, rmse

# Run evaluations
lasso_r2, lasso_rmse = evaluate(lasso, "Lasso")
ridge_r2, ridge_rmse = evaluate(ridge, "Ridge")
elastic_r2, elastic_rmse = evaluate(elastic, "Elastic Net")

# For Lasso, Ridge, ElasticNet
models = {'Lasso': lasso, 'Ridge': ridge, 'ElasticNet': elastic}

for name, model in models.items():
    print(f"{name} Coefficients:")
    for feat, coef in zip(X.columns, model.coef_):
        print(f"{feat}: {coef:.4f}")
    print("-" * 40)

```

Lasso Performance:

R<sup>2</sup>: 0.393

RMSE: 0.068

-----

Ridge Performance:

R<sup>2</sup>: -0.481

RMSE: 0.106

-----

Elastic Net Performance:

R<sup>2</sup>: 0.389

RMSE: 0.068

-----

Lasso Coefficients:

VWAP: -0.0031

Accrual Ratio: -0.0000

B/P: 0.0096

CF/P: 0.0000

Capital Expenditure: -0.0000

Cash: -0.0000

Debt/Equity: 0.0000

Depreciation: -0.0000

Dividend: -0.0000

E/P: 0.0001

EBIT: -0.0000

EBIT/P: 0.0000

EBIT/TEV: 0.0000  
 Earnings: -0.0000  
 Earnings Growth (1Y): 0.0000  
 Equity: -0.0000  
 FCF: -0.0000  
 FCF/P: 0.0000  
 Income Tax: -0.0000  
 Interest Expense: -0.0217  
 Long Liabilities: 0.0501  
 Long Term Debt: -0.0000  
 Market Cap: -0.0000  
 Minority Interest: -0.0000  
 Operating Cash Flow: -0.0000  
 Operating Expense: -0.0000  
 Operating Income: -0.0000  
 Operating Income Before Depreciation: -0.0000  
 Operating Margin: 0.0000  
 Preferred Stock: 0.0153  
 Profit Margin: 0.0000  
 R&D: -0.0000  
 R&D/Sales: -0.0000  
 ROA: -0.0000  
 ROE: -0.0000  
 S/P: -0.0000  
 SG&A: -0.0000  
 SG&A/Sales: 0.0078  
 Sales: -0.0000  
 Sales Growth (1Y): 0.0000  
 Sales Growth (2Y): -0.0000  
 Sales Growth (3Y): -0.0000  
 Sales Growth (4Y): -0.0000  
 Sales Growth (5Y): -0.0000  
 Sales Variability: -0.0000  
 Short Term Debt: 0.0000  
 TEV: -0.0000  
 Working Capital: -0.0000  
 close: -0.0000  
 dividendFactor: -0.0000  
 fiscalDint: -0.0000  
 floatShares: -0.0000  
 outstandingShares: -0.0000  
 shortInterestFloat: 0.0000  
 -----  
 Ridge Coefficients:  
 VWAP: -0.0215  
 Accrual Ratio: -0.0012  
 B/P: 0.0160  
 CF/P: -0.0225



Capital Expenditure: -0.0020  
Cash: -0.0322  
Debt/Equity: 0.0084  
Depreciation: -0.0248  
Dividend: -0.0000  
E/P: 0.0324  
EBIT: 0.0146  
EBIT/P: 0.0286  
EBIT/TEV: -0.0110  
Earnings: -0.0040  
Earnings Growth (1Y): 0.0059  
Equity: -0.0200  
FCF: 0.0026  
FCF/P: 0.0179  
Income Tax: -0.0046  
Interest Expense: -0.0544  
Long Liabilities: 0.0626  
Long Term Debt: -0.0273  
Market Cap: 0.0033  
Minority Interest: -0.0028  
Operating Cash Flow: 0.0017  
Operating Expense: 0.0041  
Operating Income: 0.0015  
Operating Income Before Depreciation: 0.0174  
Operating Margin: -0.0010  
Preferred Stock: 0.0140  
Profit Margin: -0.0012  
R&D: -0.0006  
R&D/Sales: 0.0019  
ROA: -0.0371  
ROE: -0.0041  
S/P: -0.0239  
SG&A: 0.0009  
SG&A/Sales: 0.0131  
Sales: 0.0037  
Sales Growth (1Y): 0.0020  
Sales Growth (2Y): -0.0113  
Sales Growth (3Y): -0.0064  
Sales Growth (4Y): 0.0102  
Sales Growth (5Y): 0.0036  
Sales Variability: -0.0036  
Short Term Debt: 0.0494  
TEV: 0.0250  
Working Capital: 0.0009  
close: 0.0040  
dividendFactor: 0.0019  
fiscalDint: -0.0002  
floatShares: -0.0041

outstandingShares: -0.0052  
shortInterestFloat: -0.0081

-----  
ElasticNet Coefficients:

VWAP: -0.0032  
Accrual Ratio: -0.0000  
B/P: 0.0091  
CF/P: 0.0000  
Capital Expenditure: -0.0000  
Cash: -0.0000  
Debt/Equity: 0.0000  
Depreciation: -0.0000  
Dividend: -0.0000  
E/P: 0.0007  
EBIT: -0.0000  
EBIT/P: 0.0000  
EBIT/TEV: 0.0000  
Earnings: -0.0000  
Earnings Growth (1Y): 0.0000  
Equity: -0.0000  
FCF: -0.0000  
FCF/P: 0.0000  
Income Tax: -0.0000  
Interest Expense: -0.0199  
Long Liabilities: 0.0474  
Long Term Debt: -0.0000  
Market Cap: -0.0000  
Minority Interest: -0.0000  
Operating Cash Flow: -0.0000  
Operating Expense: -0.0000  
Operating Income: -0.0000  
Operating Income Before Depreciation: -0.0000  
Operating Margin: 0.0000  
Preferred Stock: 0.0163  
Profit Margin: 0.0000  
R&D: -0.0000  
R&D/Sales: -0.0000  
ROA: -0.0007  
ROE: -0.0000  
S/P: -0.0000  
SG&A: -0.0000  
SG&A/Sales: 0.0081  
Sales: -0.0000  
Sales Growth (1Y): 0.0000  
Sales Growth (2Y): -0.0000  
Sales Growth (3Y): -0.0000  
Sales Growth (4Y): -0.0000  
Sales Growth (5Y): -0.0000

```

Sales Variability: -0.0000
Short Term Debt: 0.0000
TEV: -0.0000
Working Capital: -0.0000
close: -0.0000
dividendFactor: -0.0000
fiscalDint: -0.0000
floatShares: -0.0000
outstandingShares: -0.0000
shortInterestFloat: 0.0000
-----

```

Lasso and Elastic Net outperforms Ridge because the former two models can perform feature selection by shrinking coefficients to exactly zero when there are many features

Try the regularization models by only using features with correlation > 0.2

```

[19]: X = df[reduced_features].dropna()

y = df.loc[X.index, 'FinancialRisk']
X_scaled = scaler.fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
↳ random_state=42)

```

```

[20]: from sklearn.linear_model import LassoCV, RidgeCV, ElasticNetCV

# Lasso
lasso = LassoCV(cv=5, random_state=42)
lasso.fit(X_train, y_train)

# Ridge
ridge = RidgeCV(alphas=[0.1, 1.0, 10.0], cv=5)
ridge.fit(X_train, y_train)

# Elastic Net
elastic = ElasticNetCV(cv=5, random_state=42)
elastic.fit(X_train, y_train)

```

```

[20]: ElasticNetCV(cv=5, random_state=42)

```

```

[21]: from sklearn.metrics import r2_score, mean_squared_error
import numpy as np

def evaluate(model, X_test, y_test, name="Model"):
    preds = model.predict(X_test)
    r2 = r2_score(y_test, preds)
    rmse = np.sqrt(mean_squared_error(y_test, preds))
    print(f"{name} Performance:")
    print(f"  R²: {r2:.3f}")

```

```

print(f"  RMSE: {rmse:.3f}")
print("-" * 30)

evaluate(lasso, X_test, y_test, "Lasso")
evaluate(ridge, X_test, y_test, "Ridge")
evaluate(elastic, X_test, y_test, "Elastic Net")

# For Lasso, Ridge, ElasticNet
models = {'Lasso': lasso, 'Ridge': ridge, 'ElasticNet': elastic}

for name, model in models.items():
    print(f"{name} Coefficients:")
    for feat, coef in zip(X.columns, model.coef_):
        print(f"{feat}: {coef:.4f}")
    print("-" * 40)

```

Lasso Performance:

R<sup>2</sup>: 0.250  
 RMSE: 0.086

Ridge Performance:

R<sup>2</sup>: 0.368  
 RMSE: 0.079

Elastic Net Performance:

R<sup>2</sup>: 0.239  
 RMSE: 0.087

Lasso Coefficients:

B/P: 0.0053  
 FCF/P: 0.0000  
 Long Liabilities: 0.0264  
 Preferred Stock: 0.0226  
 Short Term Debt: 0.0000

Ridge Coefficients:

B/P: 0.0207  
 FCF/P: 0.0006  
 Long Liabilities: 0.0409  
 Preferred Stock: 0.0391  
 Short Term Debt: -0.0200

ElasticNet Coefficients:

B/P: 0.0042  
 FCF/P: 0.0000  
 Long Liabilities: 0.0256  
 Preferred Stock: 0.0217

Short Term Debt: 0.0000

-----

## 2.5 Modeling 3: non-linear models

Using Random Forest to check if non-linear models can outperform regression models (because we have a lot of variables)

```
[22]: from sklearn.ensemble import RandomForestRegressor
      # Train test split
      full_features = df.drop(columns=['FinancialRisk', 'Assets',
      ↪ 'FinancialRisk_log']) # Drop target
      X = full_features.select_dtypes(include=[np.number]).dropna() # Keep only
      ↪ numerical features
      y = df.loc[X.index, 'FinancialRisk']

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42)
```

```
[23]: # Fit model
      rf = RandomForestRegressor(n_estimators=100, random_state=42)
      rf.fit(X_train, y_train)
```

```
[23]: RandomForestRegressor(random_state=42)
```

```
[24]: # Check performance
      y_pred = rf.predict(X_test)

      r2 = r2_score(y_test, y_pred)
      rmse = np.sqrt(mean_squared_error(y_test, y_pred))

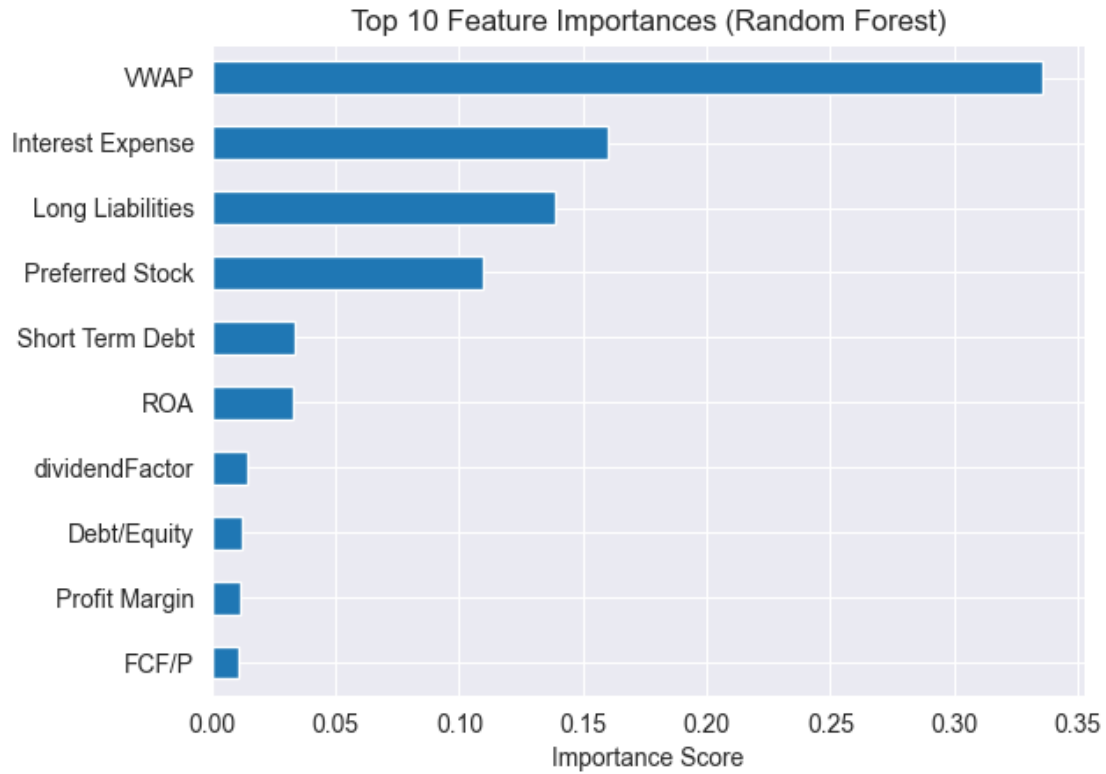
      print(f"Random Forest R²: {r2:.4f}")
      print(f"Random Forest RMSE: {rmse:.4f}")
```

Random Forest R²: -0.1234

Random Forest RMSE: 0.0927

```
[25]: # Check feature importance for rf model
      import pandas as pd
      import matplotlib.pyplot as plt

      feature_importances = pd.Series(rf.feature_importances_, index=X.columns)
      feature_importances.nlargest(10).plot(kind='barh')
      plt.title("Top 10 Feature Importances (Random Forest)")
      plt.xlabel("Importance Score")
      plt.gca().invert_yaxis()
      plt.show()
```



## # Findings

Compare performance between the 3 models

## 2.6 Base Model: Linear Regression

R<sup>2</sup>: -0.08 RMSE: 0.08

## 2.7 Regularization Model (With all features)

Ridge: R<sup>2</sup>: 0.393 RMSE: 0.068

Lasso: R<sup>2</sup>: -0.481 RMSE: 0.106

Elastic Net: R<sup>2</sup>: 0.389 RMSE: 0.068

## 2.8 Regularization Model (With features > 0.2 correlation)

Ridge: R<sup>2</sup>: 0.25 RMSE: 0.086

Lasso: R<sup>2</sup>: 0.368 RMSE: 0.079

Elastic Net: R<sup>2</sup>: 0.239 RMSE: 0.087

## 2.9 Non-linear model (Random Forest)

$R^2$ : -0.1234 RMSE: 0.0927

Best model is Ridge Model with all features included